Reinforcement Learning Exercise 3

October 15, 2019

1 Q-Learning

In this the exercise, you'll be applying grid-based Q-learning to the *Cartpole* and *LunarLander* environments.

1.1 Cartpole

Recall the Cartpole environment from Exercise 1.

Task 1 – 25 points   Implement Q-learning for the Cartpole environment in the file `qlearning.py`. Compare using a constant value of $\epsilon = 0.2$ to reducing the value of ϵ over time (use the *greedy in limit with infinite exploration* formula from the lecture). Use the following hyperparameter values: $\alpha = 0.1$ and $\gamma = 0.98$. For GLIE, aim at reaching $\epsilon = 0.1$ after 20000 episodes and round the value of a to the nearest integer.

Hint: The states in Cartpole are continuous, while grid-based methods can only be directly applied to discrete state spaces. You can overcome this issue by discretizing Cartpole states to a finite grid (use the state space limits and grid dimensions from `qlearning.py`).

Task 2 – 10 points   Use your Q-function values to calculate the optimal value function of each state. Plot the heatmap of the value function in terms of x and θ. For plotting, average the values over $\dot{x}$ and $\dot{\theta}$.

Hint: For plotting the heatmap you can use Matplotlib-pyplot: `pyplot.imshow(values_array)` or Seaborn: `seaborn.heatmap(values_array)`

Aalto University
School of Electrical
Engineering

Reinforcement Learning course staff
Intelligent Robotics Group
aalto.fi, irobotics.aalto.fi

Question 1 – 10 points   What do you think the heatmap would have looked like:
    (a) before the training?

    (b) after a single episode?

    (c) halfway through the training?

Task 3 – 5 points   Set ϵ to zero, effectively making the policy greedy w.r.t. current Q-value estimates. Run the code again,
    (a) keeping the initial estimates of the Q function at 0,

    (b) setting the initial estimates of the Q function to 50 for all states and actions.

Observe the results.

Question 2.1 – 5 points   In which case does the model perform better?

Question 2.2 – 10 points   Why is this the case?

Question 2.3 – 10 points   How does the initialization of Q values affect exploration?

1.2   Lunar lander

Figure 1: The Lunar lander environment.

The *Lunar lander* environment is shown in Figure 1. The goal is to make the lunar lander land on the ground between two flag poles. The agent receives a positive reward for moving towards the landing pad, landing, etc. A negative reward is given for firing the main engine (more fuel-efficient policies are better) and for crashing. Four actions are available: firing the left/right/main engines, or doing nothing (free fall). The observation vector consists of 6 continuous and 2 discrete values:

$$o = \begin{pmatrix} x & y & \dot{x} & \dot{y} & \theta & \dot{\theta} & c_l & c_r \end{pmatrix}^T, \tag{1}$$

where x and y are the coordinates of the lander, $\dot{x}$ and $\dot{y}$ its velocities, θ represents the rotation angle and $\dot{\theta}$ the angular velocity of the lander. Two discrete values c_l and c_r indicate whether the lander's legs are in contact with the ground (0 or 1).

Hint: For the Lunar Lander, change the environment to gym.make(LunarLander-v2).

Task 4 – 10 points Modify your code for Task 1 and try to apply it in the Lunar lander environment. Run it for 20000 episodes (which was enough for the Cartpole to learn).

Hint: The last two elements of the state vector, c_l and c_r are binary (0 or 1), so, when discretizing the state space, it is enough to use a grid dimension of two for them. This will save a substantial amount of memory.

Hint: If you're getting errors about Box2D not being installed, you have to install the box2d-py package with pip3. You may also have to additionally install Swig from your distro repositories (apt-get install/pacman -S/emerge, etc.).

Question 3.1 – 5 points Is the lander able to learn any useful behaviour?

Question 3.2 – 10 points Why/why not?

2   Submission

Your report should be submitted as a **PDF file** as RL_Exercise3_*LastName_FirstName*.pdf. The report must include:

1. **Answers to all questions** posed in the text.
2. The **training performance plots** for each of the tasks (Task 1 - fixed and GLIE, Task 3 - for both initializations, Task 4 - Lunar Lander).
3. The **heatmap** from the end of the training (Task 2).

In addition to the report, you must submit as separate files, in the same folder as the report:

1. *NumPy* file q_values.npy, which **includes the learned Q-values for Task 1 for Cartpole with GLIE**, saved when the training has finished.
2. *NumPy* file value_func.npy, which **contains the value function for the same conditions as in the previous point**.
3. Python code used to solve the exercises.

Do not attach the Q values for Lunar Lander. Please, avoid uploading unnecessary files such as self-generated folders (e.g. __pycache__, __MACOSX).

Deadline to submit your answers is **28th October 2019, 11:55 am** (in the morning before the Monday exercise session).

If you need help solving the tasks, you are welcome to visit the exercise sessions on Monday, Tuesday and Wednesday.

Good luck!

**Aalto University
School of Electrical
Engineering**

Reinforcement Learning course staff
Intelligent Robotics Group
aalto.fi, irobotics.aalto.fi